

# Storage Strategies

**Unit-7**

# Basic Concepts

- ▶ Indexing mechanisms used to speed up access to desired data.
  - ▶ E.g., author catalog in library
- ▶ **Search Key** - attribute to set of attributes used to look up records in a file.
- ▶ An **index file** consists of records (called **index entries**) of the form

search-key	pointer
------------	---------

- ▶ Index files are typically much smaller than the original file
- ▶ Two basic kinds of indices:
  - ▶ **Ordered indices:** search keys are stored in sorted order
  - ▶ **Hash indices:** search keys are distributed uniformly across “buckets” using a “hash function”.

# INDEX EVALUATION METRICS

- ▶ Access types supported efficiently. E.g.,
  - ▶ records with a specified value in the attribute
  - ▶ or records with an attribute value falling in a specified range of values.
- ▶ Access time
- ▶ Insertion time
- ▶ Deletion time
- ▶ Space overhead

# INDEXING

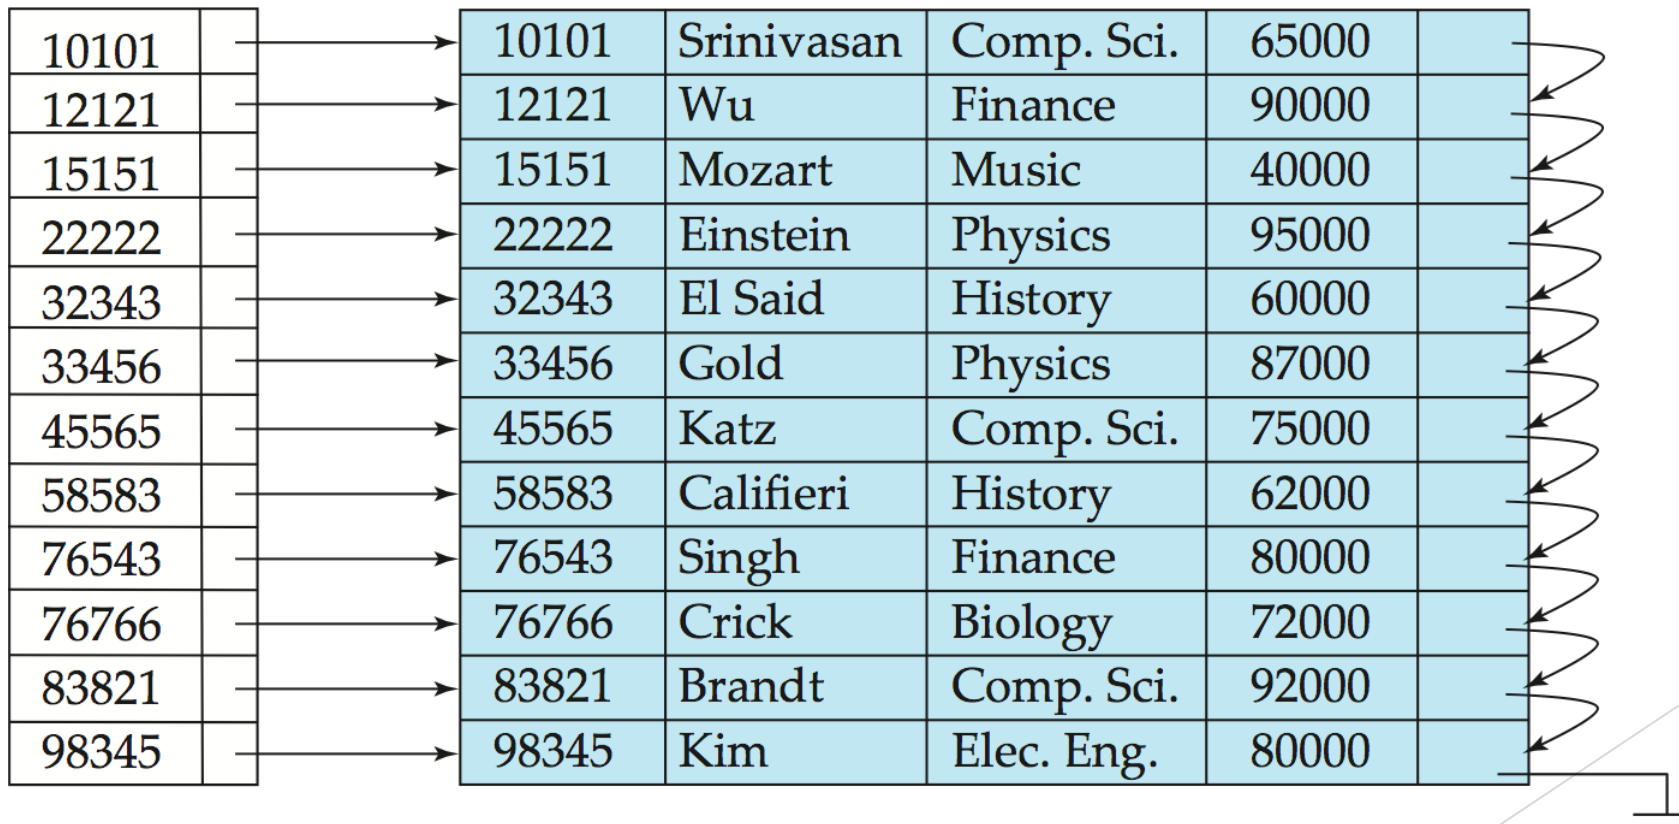
# Ordered Indices

- ▶ In an **ordered index**, index entries are stored sorted on the search key value. E.g., author catalog in library.
- ▶ **Primary index**: in a sequentially ordered file, the index whose search key specifies the sequential order of the file.
  - ▶ Also called **clustering index**
  - ▶ The search key of a primary index is usually but not necessarily the primary key.
- ▶ **Secondary index**: an index whose search key specifies an order different from the sequential order of the file. Also called **non-clustering index**.
- ▶ **Index-sequential file**: ordered sequential file with a primary index.

# Dense Index Files

- ▶ **Dense index** — Index record appears for every search-key value in the file.
- ▶ E.g. index on *ID* attribute of *instructor* relation

10101	→	10101	Srinivasan	Comp. Sci.	65000	↙
12121	→	12121	Wu	Finance	90000	↙
15151	→	15151	Mozart	Music	40000	↙
22222	→	22222	Einstein	Physics	95000	↙
32343	→	32343	El Said	History	60000	↙
33456	→	33456	Gold	Physics	87000	↙
45565	→	45565	Katz	Comp. Sci.	75000	↙
58583	→	58583	Califieri	History	62000	↙
76543	→	76543	Singh	Finance	80000	↙
76766	→	76766	Crick	Biology	72000	↙
83821	→	83821	Brandt	Comp. Sci.	92000	↙
98345	→	98345	Kim	Elec. Eng.	80000	↙



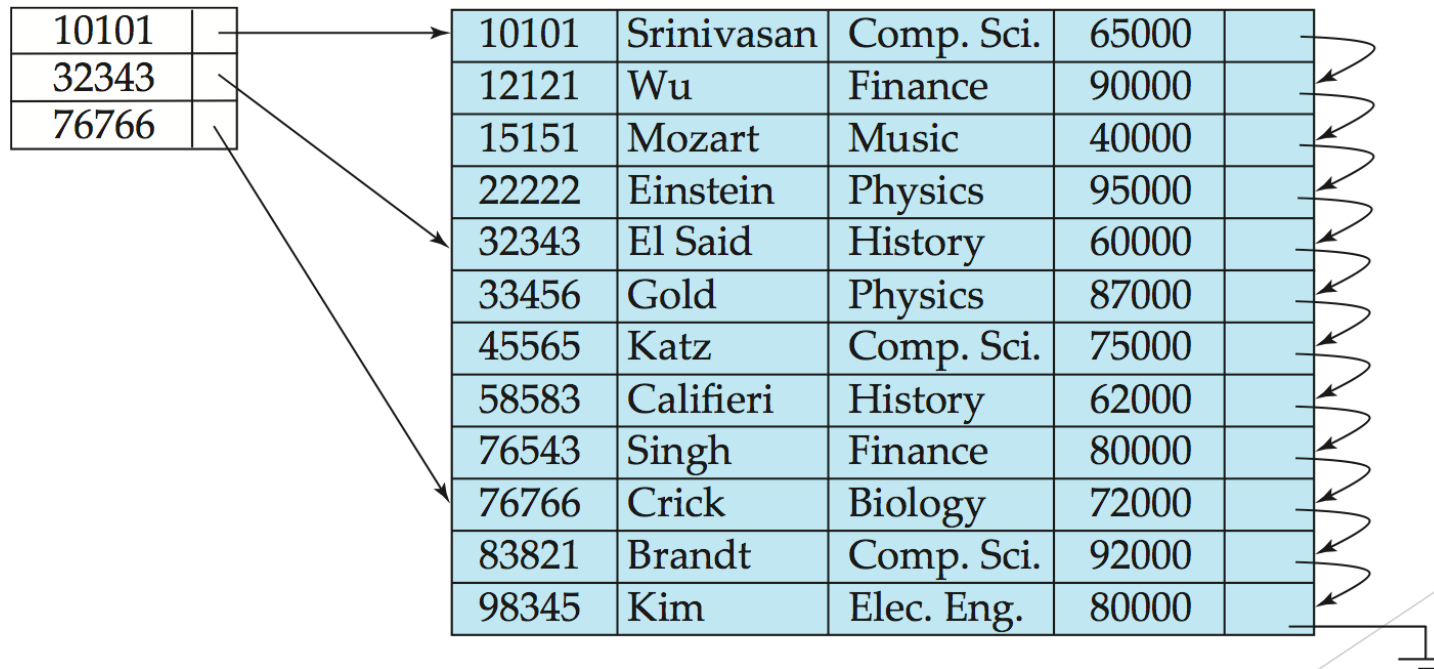
# Dense Index Files (Cont.)

- Dense index on *dept\_name*, with *instructor* file sorted on *dept\_name*

Biology		76766	Crick	Biology	72000	
Comp. Sci.		10101	Srinivasan	Comp. Sci.	65000	
Elec. Eng.		45565	Katz	Comp. Sci.	75000	
Finance		83821	Brandt	Comp. Sci.	92000	
History		98345	Kim	Elec. Eng.	80000	
Music		12121	Wu	Finance	90000	
Physics		76543	Singh	Finance	80000	
		32343	El Said	History	60000	
		58583	Califieri	History	62000	
		15151	Mozart	Music	40000	
		22222	Einstein	Physics	95000	
		33465	Gold	Physics	87000	

# Sparse Index Files


- ▶ **Sparse Index:** contains index records for only some search-key values.
  - ▶ Applicable when records are sequentially ordered on search-key
- ▶ To locate a record with search-key value  $K$  we:
  - ▶ Find index record with largest search-key value  $< K$
  - ▶ Search file sequentially starting at the record to which the index record points





# Sparse Index Files (Cont.)


- ▶ Compared to dense indices:
  - ▶ Less space and less maintenance overhead for insertions and deletions.
  - ▶ Generally slower than dense index for locating records.
- ▶ **Good tradeoff:** sparse index with an index entry for every block in file, corresponding to least search-key value in the block.



# Multilevel Index

- ▶ If primary index does not fit in memory, access becomes expensive.
- ▶ Solution: treat primary index kept on disk as a sequential file and construct a sparse index on it.
  - ▶ outer index – a sparse index of primary index
  - ▶ inner index – the primary index file
- ▶ If even outer index is too large to fit in main memory, yet another level of index can be created, and so on.
- ▶ Indices at all levels must be updated on insertion or deletion from the file.

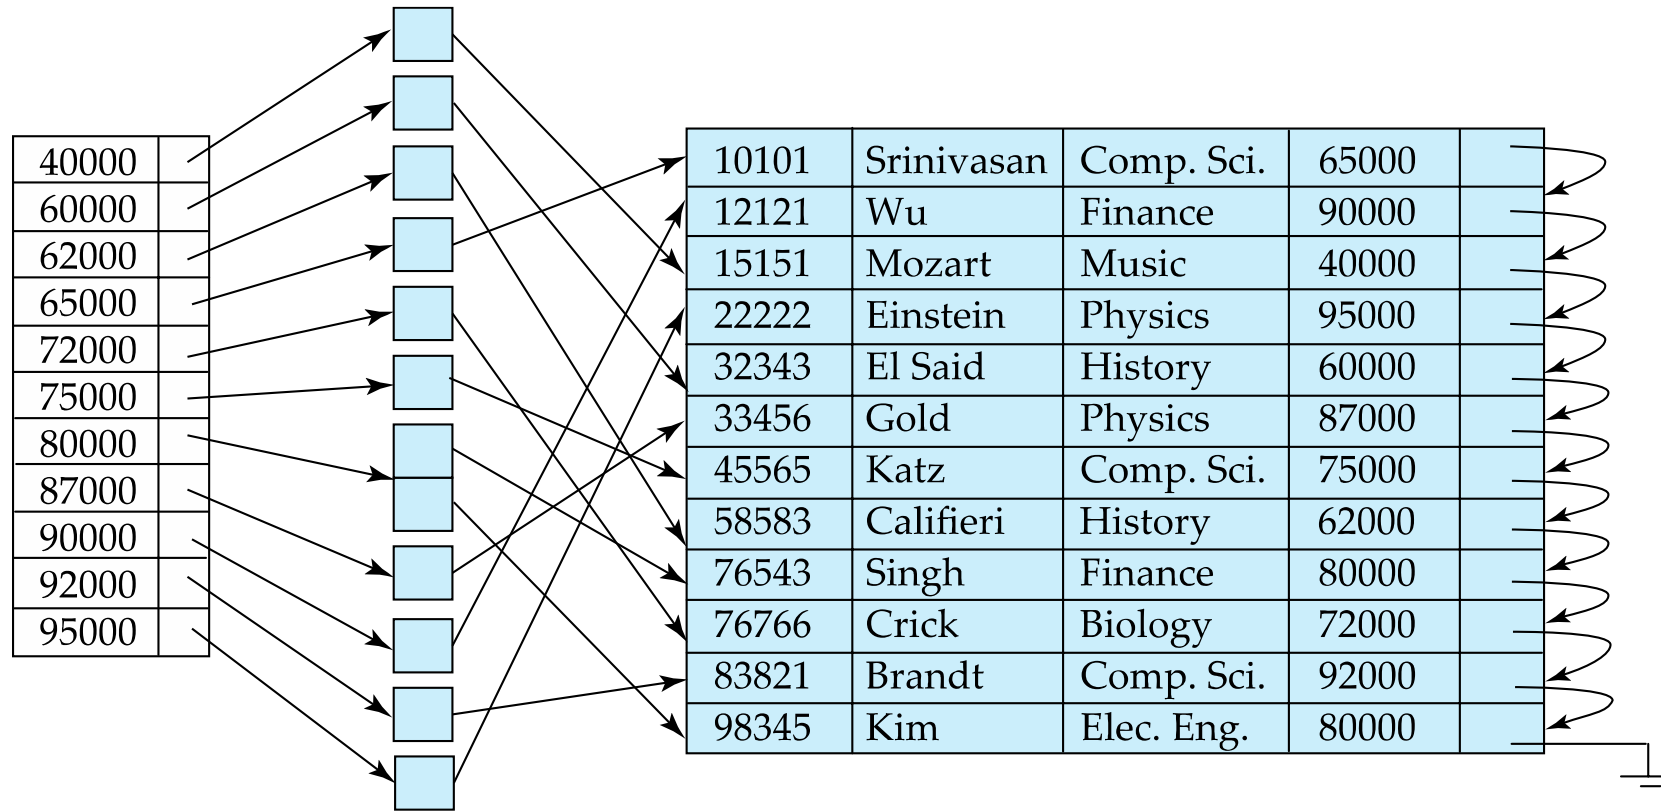
## Multilevel Index (Cont.)



# Secondary Indices

- ▶ Frequently, one wants to find all the records whose values in a certain field (which is not the search-key of the primary index) satisfy some condition.
  - ▶ Example 1: In the *instructor* relation stored sequentially by ID, we may want to find all instructors in a particular department
  - ▶ Example 2: as above, but where we want to find all instructors with a specified salary or with salary in a specified range of values
- ▶ We can have a secondary index with an index record for each search-key value

# Secondary Indices Example



**Secondary index on *salary* field of *instructor***

- ▶ Index record points to a bucket that contains pointers to all the actual records with that particular search-key value.
- ▶ Secondary indices have to be dense


# B<sup>+</sup>-Tree

# B<sup>+</sup>-Tree Index Files

A B<sup>+</sup>-tree is a rooted tree satisfying the following properties:

- ▶ All paths from root to leaf are of the same length
- ▶ Each node that is not a root or a leaf has between  $\lceil n/2 \rceil$  and  $n$  children.
- ▶ A leaf node has between  $\lceil (n-1)/2 \rceil$  and  $n-1$  values
- ▶ Special cases:
  - ▶ If the root is not a leaf, it has at least 2 children.
  - ▶ If the root is a leaf (that is, there are no other nodes in the tree), it can have between 0 and  $(n-1)$  values.

# B<sup>+</sup>-Tree Node Structure



- ▶ Typical node
  - ▶  $K_i$  are the search-key values
  - ▶  $P_i$  are pointers to children (for non-leaf nodes) or pointers to records or buckets of records (for leaf nodes).
- ▶ The search-keys in a node are ordered

$$K_1 < K_2 < K_3 < \dots < K_{n-1}$$

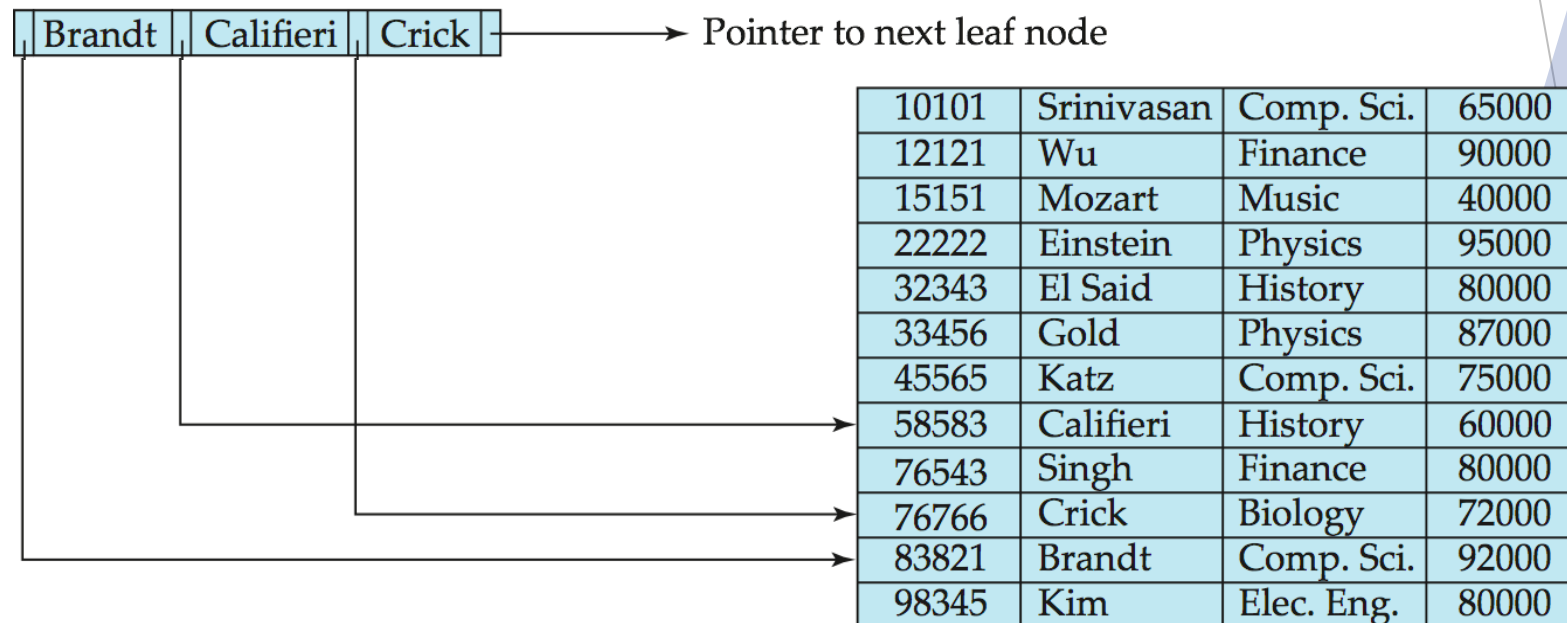
(Initially assume no duplicate keys, address duplicates later)



# Leaf Nodes in B<sup>+</sup>-Trees

## Properties of a leaf node:

- ▶ For  $i = 1, 2, \dots, n-1$ , pointer  $P_i$  points to a file record with search-key value  $K_i$ ,
- ▶ If  $L_i, L_j$  are leaf nodes and  $i < j$ ,  $L_i$ 's search-key values are less than or equal to  $L_j$ 's search-key values
- ▶  $P_n$  points to next leaf node in search-key order  
leaf node




# Non-Leaf Nodes in B<sup>+</sup>-Trees

- ▶ Non leaf nodes form a multi-level sparse index on the leaf nodes. For a non-leaf node with  $m$  pointers:
  - ▶ All the search-keys in the subtree to which  $P_1$  points are less than  $K_1$
  - ▶ For  $2 \leq i \leq n - 1$ , all the search-keys in the subtree to which  $P_i$  points have values greater than or equal to  $K_{i-1}$  and less than  $K_i$
  - ▶ All the search-keys in the subtree to which  $P_n$  points have values greater than or equal to  $K_{n-1}$

$P_1$	$K_1$	$P_2$	$\dots$	$P_{n-1}$	$K_{n-1}$	$P_n$
-------	-------	-------	---------	-----------	-----------	-------

# Example of B<sup>+</sup>-tree



B<sup>+</sup>-tree for *instructor* file ( $n = 6$ )

- ▶ Leaf nodes must have between 3 and 5 values ( $\lceil (n-1)/2 \rceil$  and  $n - 1$ , with  $n = 6$ ).
- ▶ Non-leaf nodes other than root must have between 3 and 6 children ( $\lceil n/2 \rceil$  and  $n$  with  $n = 6$ ).
- ▶ Root must have at least 2 children.

# Hashing

# Static Hashing

- ▶ A **bucket** is a unit of storage containing one or more records (a bucket is typically a disk block).
- ▶ In a **hash file organization** we obtain the bucket of a record directly from its search-key value using a **hash function**.
- ▶ Hash function  $h$  is a function from the set of all search-key values  $K$  to the set of all bucket addresses  $B$ .
- ▶ Hash function is used to locate records for access, insertion as well as deletion.
- ▶ Records with different search-key values may be mapped to the same bucket; thus entire bucket has to be searched sequentially to locate a record.

# Example of Hash File Organization

Hash file organization of *instructor* file, using *dept\_name* as key  
(See figure in next slide.)

- ▶ There are 10 buckets,
- ▶ The binary representation of the  $i$ th character is assumed to be the integer  $i$ .
- ▶ The hash function returns the sum of the binary representations of the characters modulo 10
  - ▶ E.g.  $h(\text{Music}) = 1$        $h(\text{History}) = 2$   
          $h(\text{Physics}) = 3$     $h(\text{Elec. Eng.}) = 3$

# Example of Hash File Organization

bucket 0


bucket 1

15151	Mozart	Music	40000

bucket 2

32343	El Said	History	80000
58583	Califieri	History	60000

bucket 3

22222	Einstein	Physics	95000
33456	Gold	Physics	87000
98345	Kim	Elec. Eng.	80000

bucket 4

12121	Wu	Finance	90000
76543	Singh	Finance	80000

bucket 5

76766	Crick	Biology	72000

bucket 6

10101	Srinivasan	Comp. Sci.	65000
45565	Katz	Comp. Sci.	75000
83821	Brandt	Comp. Sci.	92000

bucket 7


Hash file organization of *instructor* file, using *dept\_name* as key (see previous slide for details).

# Hash Functions

- ▶ Worst hash function maps all search-key values to the same bucket; this makes access time proportional to the number of search-key values in the file.
- ▶ An ideal hash function is **uniform**, i.e., each bucket is assigned the same number of search-key values from the set of *all* possible values.
- ▶ Ideal hash function is **random**, so each bucket will have the same number of records assigned to it irrespective of the *actual distribution* of search-key values in the file.
- ▶ Typical hash functions perform computation on the internal binary representation of the search-key.
  - ▶ For example, for a string search-key, the binary representations of all the characters in the string could be added and the sum modulo the number of buckets could be returned. .




# Handling of Bucket Overflows

- ▶ Bucket overflow can occur because of
  - ▶ Insufficient buckets
  - ▶ Skew in distribution of records. This can occur due to two reasons:
    - ▶ multiple records have same search-key value
    - ▶ chosen hash function produces non-uniform distribution of key values
- ▶ Although the probability of bucket overflow can be reduced, it cannot be eliminated; it is handled by using *overflow buckets*.

# Handling of Bucket Overflows (Cont.)

- ▶ **Overflow chaining** - the overflow buckets of a given bucket are chained together in a linked list.
- ▶ Above scheme is called **closed hashing**.
  - ▶ An alternative, called **open hashing**, which does not use overflow buckets, is not suitable for database applications.



# Deficiencies of Static Hashing

- ▶ In static hashing, function  $h$  maps search-key values to a fixed set of  $B$  of bucket addresses. Databases grow or shrink with time.
  - ▶ If initial number of buckets is too small, and file grows, performance will degrade due to too much overflows.
  - ▶ If space is allocated for anticipated growth, a significant amount of space will be wasted initially (and buckets will be underfull).
  - ▶ If database shrinks, again space will be wasted.
- ▶ One solution: periodic re-organization of the file with a new hash function
  - ▶ Expensive, disrupts normal operations
- ▶ Better solution: allow the number of buckets to be modified dynamically.

# Index Definition in SQL

- ▶ Create an index

**create index** <index-name> **on** <relation-name>  
(<attribute-list>)

E.g.: **create index** *index* **on** *branch(branch\_name)*

- ▶ Use **create unique index** to indirectly specify and enforce the condition that the search key is a candidate key.
  - ▶ Not really required if SQL **unique** integrity constraint is supported

- ▶ To drop an index

**drop index** <index-name>

- ▶ Most database systems allow specification of type of index, and clustering.

THANK YOU